

DALI MVP Vision: Lightweight Legal AI System



DALI - Development of Ambitious Leading Ideas

The NEW Venture Building Business Unit @ Siyada Tech

"Where legal innovation meets ambitious execution - because why settle for ordinary when you can build extraordinary legal AI solutions?"

Welcome to DALI, Siyada Tech's newest venture building powerhouse! We're not just another tech unit - we're the ambitious dreamers who turn "what if" into "watch this!" Our mission? To develop leading-edge legal AI solutions that make lawyers wonder how they ever lived without them.



MVP Vision: Zero-to-Hero Legal AI in Your Office

The Challenge

Building a sophisticated legal LLM/RAG system that runs on minimal hardware, scrapes the web for fresh legal data, and doesn't require a PhD in computer science to operate.

The DALI Solution

A lightweight, office-friendly legal AI system that combines the best open-source tools with zero-knowledge installation requirements. Think "legal AI for humans" - powerful enough for serious work, simple enough for your intern to set up during lunch break.

The Lightweight Tech Stack

1. LLM Engine: Ollama (The Zero-Knowledge Hero)

Why Ollama? Because installing an LLM shouldn't require a computer science degree!

- **Installation:** Literally one command: `curl -fsSL https://ollama.com/install.sh | sh`
- **Models Available:** Llama 3, Mistral, CodeLlama, and 50+ others
- **Hardware Requirements:** Runs on your laptop (4GB RAM minimum, 8GB recommended)
- **Usage:** `ollama run llama3` and you're chatting with AI

Alternative Consideration: ChatGPT OSS Alternatives - **Jan.ai:** 100% offline ChatGPT alternative - **GPT4ALL:** Local ChatGPT-style interface - **Alpaca-LoRA:** Fine-tuned Llama models

2. Web Scraping: Firecrawl (The AI-Powered Web Whisperer)

Why Firecrawl? Because traditional web scraping is so 2023!

- **AI-Powered:** Uses LLMs to understand web content structure
- **Clean Output:** Converts messy HTML to clean markdown
- **Dynamic Content:** Handles JavaScript-heavy sites
- **API-First:** Easy integration with your legal AI pipeline

```
from firecrawl import FirecrawlApp

app = FirecrawlApp(api_key="your-api-key")
result = app.scrape_url("https://legal-website.com")
# Returns clean, structured legal content
```

3. Vector Database: Chroma (The Lightweight Champion)

Why Chroma? Because you don't need a database administrator to run a database!

- **Zero Config:** Works out of the box
- **Lightweight:** Perfect for office deployments

- **Python Native:** Integrates seamlessly with your Python stack

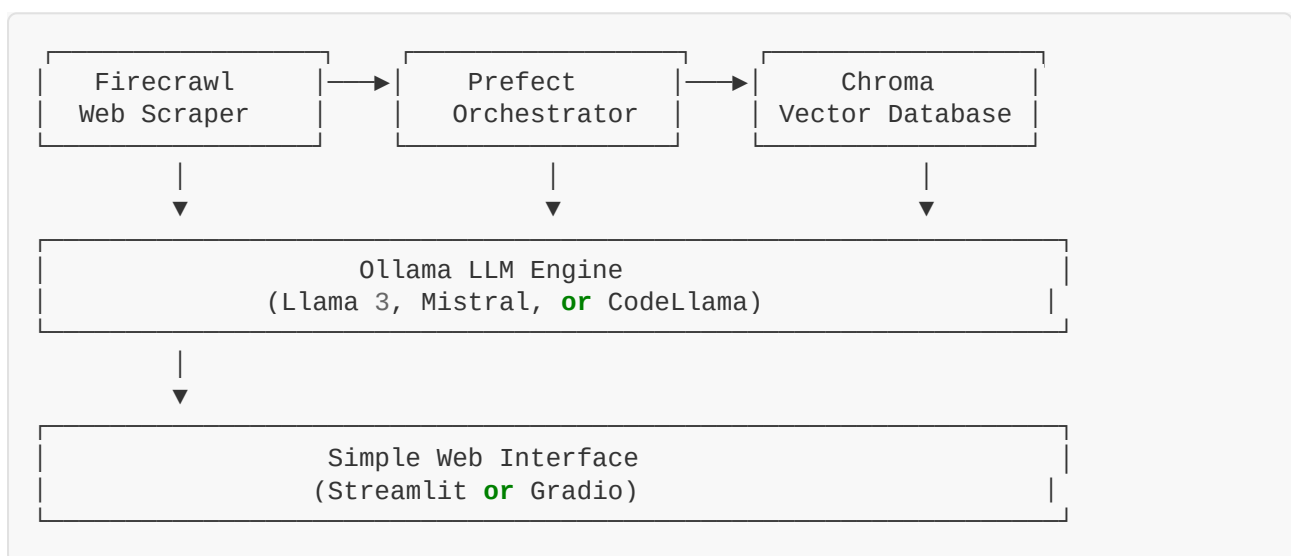
4. Workflow Orchestration: Prefect (The Modern Choice)

Why Prefect? Because Airflow is overkill for an MVP!

- **Pythonic:** Write workflows in pure Python
- **Local First:** Runs on your laptop
- **Beautiful UI:** Actually enjoyable to use



MVP Architecture



Minimum Hardware Requirements

The "Intern's Laptop" Spec

- **CPU:** Any modern processor (Intel i5 or AMD Ryzen 5 equivalent)
- **RAM:** 8GB minimum (16GB for comfort)
- **Storage:** 50GB free space (for models and data)
- **GPU:** Optional (CPU-only works fine for MVP)
- **OS:** Windows, macOS, or Linux

The "Partner's Workstation" Spec (Recommended)

- **CPU:** Intel i7 or AMD Ryzen 7
- **RAM:** 16GB+
- **Storage:** 100GB+ SSD
- **GPU:** NVIDIA RTX 4060 or better (for faster inference)

What's Happening in Hugging Face Land?

Latest Developments (2025)

1. **Vision Language Models:** Better, faster, stronger VLMs
2. **Transformers 4.55+:** Enhanced FlashAttention-2 support
3. **300+ Model Architectures:** Average of 3 new architectures monthly
4. **SmolaAgents:** New lightweight agentic library
5. **1M+ Model Checkpoints:** The largest model repository

Legal-Specific Models Worth Watching

- **Legal-BERT:** Fine-tuned for legal text classification
- **SaulLM-7B:** First LLM designed explicitly for legal text
- **Legal Document Embeddings:** Specialized sentence transformers

MVP Implementation Roadmap

Phase 1: Foundation (Week 1-2)

1. **Install Ollama:** `curl -fsSL https://ollama.com/install.sh | sh`
2. **Download Llama 3:** `ollama pull llama3`
3. **Set up Firecrawl:** API key and basic scraping tests
4. **Install Chroma:** `pip install chromadb`

Phase 2: Integration (Week 3-4)

1. **Build scraping pipeline** with Firecrawl
2. **Create vector embeddings** with Chroma
3. **Implement RAG pipeline** with Ollama
4. **Basic web interface** with Streamlit

Phase 3: Enhancement (Week 5-6)

1. **Add Prefect orchestration**
2. **Implement data refresh workflows**
3. **Add legal document classification**
4. **User authentication and access control**



The DALI Advantage

Why This Stack Rocks

1. **Zero DevOps:** Everything runs locally
2. **Zero Licensing Costs:** All open-source
3. **Zero PhD Required:** Simple installation and operation
4. **Maximum Flexibility:** Easy to customize and extend

What Makes It Legal-Ready

1. **Data Privacy:** Everything stays in your office
2. **Compliance:** No external API calls for sensitive data
3. **Customization:** Fine-tune models on your legal documents
4. **Scalability:** Start small, grow as needed



Getting Started: The 30-Minute Challenge

Can you get a basic legal AI system running in 30 minutes? With DALI's stack, absolutely!

```
# Install Ollama (1 minute)
curl -fsSL https://ollama.com/install.sh | sh

# Download Llama 3 (5 minutes)
ollama pull llama3

# Install Python dependencies (2 minutes)
pip install chromadb streamlit firecrawl-py prefect

# Clone DALI starter template (1 minute)
git clone https://github.com/siyada-tech/dali-legal-ai-mvp

# Run the MVP (1 minute)
cd dali-legal-ai-mvp
streamlit run app.py

# Start scraping legal data (20 minutes of coffee while it works)
python scrape_legal_data.py
```



Future Enhancements

Phase 2 Features

- **Multi-modal support:** Handle PDFs, images, and audio
- **Advanced RAG:** Implement graph-based retrieval
- **Legal entity extraction:** Identify parties, dates, and legal concepts
- **Document generation:** Create legal drafts and summaries

Phase 3 Features

- **Multi-language support:** Arabic legal documents
 - **Collaborative features:** Team-based legal research
 - **Advanced analytics:** Legal trend analysis and insights
 - **Mobile app:** Legal AI in your pocket
-

DALI - Where ambitious legal ideas become reality, one line of code at a time!

Contact the DALI Team: - Email: dali@siyada.tech - Slack: #dali-ventures - Office: The corner with the best coffee machine

"Remember: We're not just building software, we're building the future of legal practice. And we're having fun doing it!" - The DALI Team